

Discriminant and Classification Analysis of Prostate Cancer Prognostic Measurements

STAT 419: Multivariate Analysis — Spring 2026 Course Project

Curtis Bradley Ezli Cruz Andrea Pareja
Andres Rocha Jayasinha

A Introduction / Background

A team of urologists set out to study how prostate-specific antigen (PSA) and a collection of clinical measurements relate to the progression of disease in men with advanced prostate cancer. The data set summarizes seven quantitative measurements taken on each of 97 such patients, along with a categorical variable indicating the severity of the disease. The goal of this project is to determine whether those seven measurements, considered jointly, can distinguish patients according to the severity of the seminal vesicle invasion, and to identify which measurements contribute most to that distinction.

The grouping variable for all analyses is **DIAGN**, which indicates the severity of the seminal vesicle invasion. This variable takes on three levels low, moderate, and severe and is nearly balanced with 32 low observations, 33 moderate observations, and 32 severe observations.

The remaining seven variables are all quantitative and will serve as the predictors:

- **PSA** - serum prostate-specific antigen level, in mg/ml. A widely used biomarker for prostate disease; elevated values are associated with more advanced disease.
- **Volume** - volume of the prostate cancer, in cubic centimeters (cm³). A larger tumor volume generally signals a more advanced condition.
- **Weight** - weight of the prostate gland, in grams (g).
- **Age** - age of the patient, in years,
- **BPH** - amount of benign prostatic hyperplasia, in square centimeters (cm²). A measure of non-cancerous enlargement of the prostate.
- **Capsule** - degree of capsular penetration, i.e. the extent to which the tumor has pushed into or through the prostatic capsule.
- **Gleason** - a pathologically determined qualitative grade of the disease. Higher scores indicate a worse prognosis; in this sample the score takes only the integer values 6,7,8.

The analysis using these features went in three stages. First we examined the distribution of each predictor through histograms and summary statistics (Section B). We then checked for multicollinearity amongst our features, and fitted linear discriminants to find the combination of variables that best separate the three severity groups while also testing these combinations for significance (TODO). Finally, using the most influential variables, we built a linear classification functions and applied them to classify patients, and evaluate how accurately the resulting rule determines the true severity groups (TODO).

B Graphs and Summary Statistics

Before analysis, we will see the nature of each variable in our dataset. Below is a summary statistics table for each variable.

```
knitr::kable(data.frame(
  Variable = QUANT_VARS,
  Min      = sapply(cancer[QUANT_VARS], min, na.rm = TRUE),
  Q1       = sapply(cancer[QUANT_VARS], quantile, probs = 0.25, na.rm =
    ↪ TRUE),
  Mean     = sapply(cancer[QUANT_VARS], mean, na.rm = TRUE),
  Median   = sapply(cancer[QUANT_VARS], median, na.rm = TRUE),
  Q3       = sapply(cancer[QUANT_VARS], quantile, probs = 0.75, na.rm =
    ↪ TRUE),
  Max      = sapply(cancer[QUANT_VARS], max, na.rm = TRUE),
  SD       = sapply(cancer[QUANT_VARS], sd, na.rm = TRUE),
  row.names = NULL
), digits = 4)
```

Table 1

Variable	Min	Q1	Mean	Median	Q3	Max	SD
PSA	0.6510	5.6410	23.7301	13.3300	21.3280	265.0720	40.7829
Volume	0.2592	1.6653	6.9987	4.2631	8.4149	45.6042	7.8809
Weight	10.6970	29.3710	45.4914	37.3380	48.4240	450.3390	45.7051
Age	41.0000	60.0000	63.8660	65.0000	68.0000	79.0000	7.4451
BPH	0.0000	0.0000	2.5347	1.3499	4.7588	10.2779	3.0312
Capsule	0.0000	0.0000	2.2454	0.4493	3.2544	18.1741	3.7833
Gleason	6.0000	6.0000	6.8763	7.0000	7.0000	8.0000	0.7396

Below we can also so histograms of each variable.

```

par(mfrow=c(3,3))

for (var in QUANT_VARS) {
  hist(cancer[[var]],
      main = paste("Histogram of", var),
      breaks = 15,
      xlab = var)
}

```




Figure 1: Histograms for all quantitative variables

Based on the histogram it looks like the variables **PSA**, **Volume**, **Weight**, **BPH**, and **Capsule** are non normal. We can try to apply transformations on them to see if we can make them look more normal.

```

par(mfrow=c(2,3))

TRANSFORMED_VARS <- c('PSA', 'Volume', 'Weight', 'BPH', 'Capsule')
LOG_VARS <- c('PSA', 'Volume', 'Weight')
LOG1P_VARS <- setdiff(TRANSFORMED_VARS, LOG_VARS)

for (var in LOG_VARS) {
  hist(log(cancer[var][,]), main=paste("Histogram of LOG(", var, ")",
    ↪ sep=''), breaks=15, xlab=paste('log(',var,')',sep=''))
}

for (var in LOG1P_VARS) {
  hist(log1p(cancer[var][,]), main=paste("Histogram of LOG(", var, " + 1)",
    ↪ sep=''), breaks=15, xlab=paste('log(',var,' + 1)',sep=''))
}

```




Figure 2: Histograms for all transformed variables

i Note

BPH and Capsule are transformed with a $\log + 1$ transformation because they contain 0 values. See Table 1

Looking at Figure 2, PSA, Volume, and Weight all appear to be roughly normal, however Capsule and BPH do not. Because Discriminant Analysis has a normality assumption we drop them from our data and only use $\log(\text{Volume})$, $\log(\text{PSA})$, $\log(\text{Weight})$, Age and Gleason for our analysis.

Below is a scatterplot matrix of the transformed dataset.

```

cancer.mod <- cbind(cancer[setdiff(colnames(cancer),
  ↪ TRANSFORMED_VARS)], log(cancer[LOG_VARS]))

pairs(cancer.mod,
      main = "Scatterplot Matrix of Quantitative Variables",
      pch = 19, cex = 0.5,
      col = c("steelblue", "darkorange",
  ↪ "firebrick")[cancer.mod$DIAGN])

```




Figure 3: Scatterplot matrix of transformed dataset

B.1 Outlier Identification

After looking at the data we can see some trends and possibly some outliers. Of note there is a point that appears to have a higher weight than the others. We can also check for outliers by performing a slightly more statistical method. Assuming that we have an MVN already, we can see how far away each dot is from the center, and the probability of seeing an observation that extreme.

```
md <- mahalanobis(  
  x = cancer.mod[, -1],  
  center = colMeans(cancer.mod[, -1], na.rm = TRUE),  
  cov = cov(cancer.mod[, -1], use = "complete.obs")  
) # Compute Mahalanobis distance for every point to the data mean.  
  
sig_level <- 0.05 / nrow(cancer[, -1]) # Bonferroni correction  
  
observation_p <- pchisq(md, df = ncol(cancer.mod[, -1])) # Compute the p value  
  ↪ of seeing the value given MVN  
  
knitr::kable(cbind(cancer, observation_p)[which(observation_p < sig_level |  
  ↪ observation_p > (1 - sig_level)),]) # Select all points more extreme than  
  ↪ the significance level
```

Table 2

	DIAGN	PSA	Vol- ume	Weight	Age	BPH	Cap- sule	Glea- son	observa- tion_p
41	Moder- ate	7.463	1.1972	450.339	65	5.4739	0	6	0.9999931

Based on the results we can see that 41 appears to be different from the dataset. Looking at the output of Table 2 and comparing it to Table 1, all of the variables seem to be within normal range, with the exception of **Weight**. The weight of 450.339 grams is significantly higher than the normal range. In fact, it is 8.858 standard deviations away from the mean weight. For this reason we excluded this datapoint from our analysis.

```
cancer.mod <- cancer.mod[-41,]
```

We then checked our MVN assumptions.

B.2 MVN Assumptions

```
par(mfrow=c(2,2))

check.mvnorm.plot(cancer.mod[, -1])
mtext("All Values", side=3)
check.mvnorm.plot(cancer.mod[cancer.mod$DIAGN == "Low", , -1])
mtext("DIAGN=Low", side=3)
check.mvnorm.plot(cancer.mod[cancer.mod$DIAGN == "Moderate", , -1])
mtext("DIAGN=Moderate", side=3)
check.mvnorm.plot(cancer.mod[cancer.mod$DIAGN == "Severe", , -1])
mtext("DIAGN=Severe", side=3)
```




Figure 4: MVN Q-Q Plots

Table 3

```
cov_mats <- list(
  Low = cov(cancer.mod[cancer.mod$DIAGN == "Low", -1], use = "complete.obs"),
  Moderate = cov(cancer.mod[cancer.mod$DIAGN == "Moderate", -1], use =
    ↪ "complete.obs"),
  Severe = cov(cancer.mod[cancer.mod$DIAGN == "Severe", -1], use =
    ↪ "complete.obs")
)

n_obs <- unname(as.vector(table(cancer.mod$DIAGN)[c("Low", "Moderate",
  ↪ "Severe")]))

Box.M(cov_mats, n_obs)
```

```
$u
[1] 73.14333
```

```
$p.value.u
[1] 1.819254e-05
```

```
$chi.df
[1] 30
```

These plots looks pretty much like they're all MVN. The last check is to make sure they all have the same covariance matrices. To do this we use Box's M test.

Caution

We actually failed Box's M test. However we still decided to proceed with discriminant analysis as it looks close enough.

Based on the output it appears all groups follow a multivariate normal distribution. This allows us to move on to discriminant analysis.

C Discriminant and Classification Analysis

C.1 Checking Multicollinearity

```
knitr::kable(cor(cancer.mod[, -1]))
```

Table 4

	Age	Gleason	PSA	Volume	Weight
Age	1.0000000	0.2294805	0.1707274	0.2278361	0.3470233
Gleason	0.2294805	1.0000000	0.5384069	0.5068264	0.2226661
PSA	0.1707274	0.5384069	1.0000000	0.7340158	0.4353842
Volume	0.2278361	0.5068264	0.7340158	1.0000000	0.2857982
Weight	0.3470233	0.2226661	0.4353842	0.2857982	1.0000000

To assess for multicollinearity, we analyze the correlation coefficients for each pair of quantitative variables. Our scatterplot matrix shows that there is no prominent linear relationship between any two variables. We can then look at the calculated correlation coefficients table. Using a correlation threshold of $|r| \geq 0.80$, none of our variables have high correlation. We will not flag any of our variables for multicollinearity. We will keep all of our 5 quantitative variables.

C.2 Discriminant Functions and Variable Importance

```
da <- discrim(cancer.mod[, -1], cancer.mod[, 1])

n_ld <- ncol(da$a.stand)

std_coef <- da$a.stand
rownames(std_coef) <- colnames(cancer.mod[, -1])
colnames(std_coef) <- paste0("LD", seq_len(n_ld))

knitr::kable(std_coef)
```

Table 5

	LD1	LD2
Age	-0.0478356	-0.1523041
Gleason	-0.1538214	-0.1146393
PSA	-0.5899614	0.2450616
Volume	-0.0681763	-0.4576420
Weight	-0.0426631	0.2893474

A linear discriminant analysis was performed to observe possible separation between the groups of severity of seminal vesicle invasion and what variables contribute the most in that separation. The discriminant analysis resulted in 2 discriminant functions:

$$LD1 = -0.0478 \text{ Age} - 0.1538 \text{ Gleason} - 0.5900 \log(\text{PSA}) - 0.0682 \log(\text{Volume}) - 0.0427 \log(\text{Weight})$$

$$LD2 = -0.1523 \text{ Age} - 0.1146 \text{ Gleason} + 0.2451 \log(\text{PSA}) - 0.4576 \log(\text{Volume}) + 0.2893 \log(\text{Weight})$$

.

When looking at the importance of each variable in the first discriminant function, we get:

```
ld1_order <- names(sort(abs(std_coef[, 1]), decreasing = TRUE))

knitr::kable(
  data.frame(
    Variable = ld1_order,
    LD1 = std_coef[ld1_order, 1],
    row.names = NULL
  ),
  digits = 4,
  caption = "LD1 coefficients sorted by importance"
)
```

Table 6: LD1 coefficients sorted by importance

Variable	LD1
PSA	-0.5900
Gleason	-0.1538
Volume	-0.0682
Age	-0.0478
Weight	-0.0427

The table of the standardized coefficients allows us to see that PSA seems to strongly contribute in group separation of severity of seminal vesicle invasion. Then Gleason seems to have moderate contributions while Volume, Age, and Weight seem to have the weakest contribution.

C.3 Test of Significance of the Discriminant Functions

To see which of linear discriminant functions are significant in group separation, we will test:

Test 1 $H_0: \alpha_1 = \alpha_2 = 0$

H_a : at least one $\alpha \neq 0$

Test 2 $H_0: \alpha_2 = 0$ $H_a: \alpha_2 \neq 0$

Since we have two discriminant functions, there will be 2 tests performed (listed above), we must adjust our level of significance to $\frac{0.05}{2} = 0.025$. The results are:

```
knitr::kable(discr.sig(cancer.mod[, -1], cancer.mod[, 1]))
```

Table 7

	Lambda	V	p.values
LD1	0.2456491	127.750456	0.0000000
LD2	0.9476667	4.891468	0.2986159

LD1: With a p-value of 0.00 being less than our adjusted significance level of 0.025, we have enough evidence that the first linear discriminant function works well in group separation of severity of seminal vesicle invasion.

LD2: With a p-value of 0.2986 being larger than our adjusted significance level of 0.025, we don't have enough evidence that the second discriminant function is useful in group separation of severity of seminal vesicle invasion.

C.4 Significance of Each Variable, Adjusted for the Others

After seeing the standardized coefficients of the first discriminant function, the variables PSA and Capsule might be important in group separation. Here are the formal partial F tests for each variable:

H_0 : Variable (y) does not contribute to group separation of severity of seminal vesicle invasion, adjusting for the other variables.

H_a : Variable (y) does contribute to group separation of severity of seminal vesicle invasion, adjusting for the other variables.

Since there are 7 tests being performed, we must adjust our significance level to $\alpha^* = \frac{.05}{7} = 0.0071$.

```
knitr::kable(partial.F(cancer.mod[, -1], cancer.mod[, 1]))
```

Table 8

	Lambda	F.stat	p.value
PSA	0.5688174	33.7324827	0.0000000
Gleason	0.9625959	1.7291599	0.1833405
Volume	0.9672944	1.5046102	0.2276985
Weight	0.9855378	0.6530123	0.5229511
Age	0.9934851	0.2918139	0.7476194

From our results, PSA significantly contributes the most in group separation of severity of seminal vesicle invasion, adjusting for the other variables.

All other variables do not individually contribute significantly in group separation of severity of seminal vesicle invasion, after adjusting for the other variables.

C.5 Plot of the First Two Discriminant Functions

Below is a plot of the two linear discriminant functions.

```
discr.plot(cancer.mod[, -1], cancer.mod[, 1], title='Discriminant Plot for
↪ Diagnosis Groups', leg = levels(cancer.mod[, 1]))
```




Figure 5: Discriminant Plots

From this plot, it is visually evident the first discriminant function (LD1) captures the group separation well of the three groups of low, moderate, and severe seminal vesicle invasion. The severe severity of seminal vesicle invasion is on the left while the moderate level is on the right. Low severity of seminal vesicle invasion seems grouped in between the moderate and severe levels.

The second discriminant function (LD2) does not capture the separation nature of the data as projections of the points on the line are clumped together. This further proves that LD1 is the better discriminant function that works well in group separation.

C.6 Classification Analysis

For classification analysis we first select the top 4 most important variables.

— TODO —

Linear Classification Functions

D Summary

The goal of this project was to determine whether seven clinical measurements could effectively distinguish prostate cancer patients according to the severity of seminal vesicle invasion. Based on our analysis, the answer is yes. Among the measured variables, PSA was the most influential in separating the groups, as evidenced by its significance in the discriminant analysis and its large contribution to the first linear discriminant function. In addition, plotting the observations with respect to LD1 and LD2 revealed strong separation among the severity groups, particularly along LD1.

This conclusion is further supported by the classification results. Using the linear classification functions, the apparent correct classification rate (ACCR) was 0.8646, meaning that 83 of the 96 observations were correctly classified. The confusion matrix also indicates that the model avoided the most severe types of misclassification: no patients in the low-severity group were classified as severe, and no patients in the severe group were classified as low. Instead, the misclassifications were limited to adjacent severity categories, such as low-severity observations being classified as moderate. Overall, these results suggest that the seven clinical measurements, particularly PSA, provide meaningful discrimination among levels of seminal vesicle invasion severity.

One thing in this analysis did stand out, mainly that no variables were collinear enough to drop in Section C.1, even though we expected PSA, Volume, and Capsule might have been redundant; their correlations were moderate, but never reached our thresholds.

If we were to repeat this project, we would make two key improvements. First, we would explore alternative data transformations, either in place of or in combination with log transformations, such as the Box–Cox transformation. This would be particularly useful for variables such as BPM and Capsule, which remain skewed even after applying a log transformation. Second, we would implement leave-one-out cross-validation to estimate the classification error rate. This would allow us to assess whether the selected transformations and resulting classification model generalize well to new observations, rather than overfitting to the particular dataset used in this analysis.

Finally, one variable we wish the data would have included is a direct measure of tumor stage (such as a TNM clinical stage) or a PSA reading taken over time rather than at a single point. More generally some variable(s) that would give information about the trajectory of the disease for a given patient would help in getting cleaner separation for the Moderate group, which was the hardest to classify with the measurements available.